

# **Project Documentation: Arabic-English Translation using MarianMT**

## **Project Title**

Arabic-English Neural Machine Translation using MarianMT

---

## **Team Member Names:**

- Yassmin ahmed ibrahim
- Maram essam ahmed
- Youssef tarek aboelfotouh
- Warda Khaled dahy
- Mariam ahmed Mostafa

Youssef Mohamed fathy

## Overview

This project focuses on building a neural machine translation (NMT) system that translates English text into Arabic using the pre-trained MarianMT model from Hugging Face. The system is trained and evaluated using a parallel corpus of English-Arabic sentence pairs. The implementation is based on the `transformers` library, and the training pipeline leverages Hugging Face's `Seq2SeqTrainer`.

---

## 1. Dataset Information

**Dataset Source :** Kaggle

- **Name:** Arabic to English Translation Sentences

- **File Used:** ara\_eng.txt
- **Columns:** english, arabic
- **Format:** Tab-separated values (TSV)

## **Dataset Preprocessing**

- Removed punctuation and unnecessary characters.
- Lowercased English text.
- Cleaned Arabic text using Unicode ranges for Arabic.
- Filtered out sentence pairs with less than 3 words or less than 5 characters.
- Removed duplicate English entries.

## **Data Split**

- **Training Set:** 80%
- **Testing Set:** 20%

## **Preprocessing Steps**

### **.English Text:**

Converted text to lowercase.

Removed brackets and their content: (...), [...].

Removed unnecessary characters while keeping useful punctuation: ., ,, !, ?.

Normalized characters using `unicodedata.normalize()` to standardize accented letters.

### **.Arabic Text:**

Removed all non-Arabic characters using Unicode range: `\u0600-\u06FF`.

Removed diacritics and decorative marks, such as: ﷻ.

Removed digits and non-alphabetic symbols.

---

## **2. Model Information**

### **Pre-trained Model**

- **Model Name:** `Helsinki-NLP/opus-mt-tc-big-en-ar`
- **Library:** `transformers` by Hugging Face

- **Architecture:** MarianMT (based on Transformer architecture)

## Tokenizer

- **Type:** SentencePiece
- **Tokenizer:** `MarianTokenizer` from Hugging Face

## Input Length

- **Maximum Sequence Length:** 128 tokens
- 

## 3. Training Details

### Preprocessing

- Tokenization applied to both English input and Arabic target.
- Padding and truncation used for fixed-length inputs (max length = 128).

### Training Configuration

- **Epochs:** 10
- **Training Batch Size:** 16
- **Evaluation Batch Size:** 16
- **Output Directory:** `./results`
- **Logging Directory:** `./logs`

- **Logging Steps:** 10
- **Save Steps:** 1000
- **Evaluation Strategy:** Custom bertscore callback
- **Prediction:** Enabled with `predict_with_generate=True`

### Training Loss (last few steps):

- Decreased from ~6.5 to ~0.2 over 10 epochs
  - Indicates effective learning and convergence
- 

## 4. Evaluation Metrics

### Metric Used: BERTScore

- **Library:** `bert_score` (from HuggingFace)
  - **Language Models Used:** Pre-trained multilingual BERT models (`bert-base-multilingual-cased`)
  - **Method:**
    - Computes similarity between reference and generated translation using contextual embeddings
    - Captures semantic meaning better than lexical match-based metrics like BLEU
  - **Evaluation Frequency:**
    - Evaluated on 100 random test samples at the end of each epoch
    - Provides F1 score for precision-recall balance
  - **Why BERTScore?**
    - More reliable for low-resource language pairs
    - Better reflects human judgment of translation quality
-

## 5. Model Limitations

- **Limited Domain Generalization:**
  - Trained on a small parallel corpus; may not generalize to other domains (e.g., medical, legal).
- **Token Truncation:**
  - Long sentences may be truncated, potentially affecting translation accuracy.
- **Vocabulary Limitation:**
  - Pre-trained tokenizer might not handle rare or domain-specific tokens well.
- **Cultural and Contextual Nuances:**
  - May miss idiomatic expressions or cultural references that require contextual understanding.
- **BERTScore Reliability:**
  - While more semantically aware than BLEU, still limited by the capabilities of the underlying BERT model
  - Sensitive to language-specific pretraining and sentence structure nuances.

---

## 6.Parameters Model

1. `model_name`
  - **Type:** `str`

- **Description:** The name of the pre-trained Marian model used for English-to-Arabic translation. It is used to load the model and tokenizer.

## 2. `max_length`

- **Type:** `int`
- **Description:** The maximum token length for the input and output sequences during tokenization.

## 3. `train_dataset & test_dataset`

- **Type:** `Dataset`
- **Description:** The training and testing datasets that are used for training and evaluation. These are instances of the `Dataset` class from the `datasets` library.

## 4. `training_args`

- **Type:** `Seq2SeqTrainingArguments`
- **Description:** This contains all the hyperparameters required for training the model. It specifies batch size, number of epochs, logging settings, save steps, etc.
- **Important fields:**
  - `per_device_train_batch_size`: The batch size for training.
  - `num_train_epochs`: The number of epochs for training.
  - `logging_dir`: Directory where logs will be saved.
  - `save_steps`: Frequency at which the model is saved.

## 5. `generation_config`

- **Type:** `GenerationConfig`
- **Description:** Configuration for controlling the generation of text by the model.
- **Important fields:**
  - `max_length`: Maximum length of the generated sequence.
  - `num_beams`: Number of beams for beam search (controls the search size during translation).
  - `length_penalty`: Penalty for longer translations.
  - `early_stopping`: Whether to stop generating early if the sequence is finished.



- `bad_words_ids`: A list of words to avoid in the generated text.

#### 6. `batch_size`

- **Type:** `int`
- **Description:** The number of samples in each batch during translation.

#### 7. `inputs` (in `preprocess` function)

- **Type:** `dict`
- **Description:** This is the tokenized input for the English text. It is a dictionary containing keys like `input_ids` and `attention_mask`, which are used by the model to process the input.

#### 8. `labels` (in `preprocess` function)

- **Type:** `dict`
- **Description:** The tokenized version of the Arabic text (target). It also contains `input_ids` which represent the tokens for the target language.

#### 9. `decoded_preds` and `decoded_labels`

- **Type:** `list of str`
- **Description:** The translated output (`decoded_preds`) and the reference translation (`decoded_labels`). These are lists of strings containing the decoded translations.

#### 10. `P, R, F1`

- **Type:** `Tensor`
- **Description:** The Precision (P), Recall (R), and F1 scores computed using the BERTScore metric to evaluate the translation quality.

#### 11. `preds` and `labels` (in `compute_metrics`)

- **Type:** `list of int`
- **Description:** These are the predicted and true label sequences (after tokenization) used for computing metrics like BERTScore.

#### 12. `texts` (in `generate_translation`)

- **Type:** `list of str`
- **Description:** The English text to be translated. This is passed in batches to the model for translation.

#### 13. `input_ids` and `attention_mask` (in `generate_translation`)

- **Type:** `Tensor`
  - **Description:** These are the inputs to the model during generation. `input_ids` are the tokenized representation of the input texts, and `attention_mask` is used to indicate which tokens should be attended to (ignores padding tokens).
14. **interface (in Gradio UI)**
- **Type:** `gr.Interface`
  - **Description:** This is the Gradio interface that takes an English sentence as input and displays the predicted Arabic translation as output. It uses the `generate_translation` function as its backend processing logic.

## 7. Future Enhancements

- **Data Augmentation:**
  - Use back-translation or synonym replacement to enrich dataset.
- **Fine-Tuning:**
  - Use a larger or more domain-specific dataset.
- **UI Integration:**
  - Integrate with a GUI (e.g., Gradio or Streamlit) for easy testing and deployment.
- **Deploy API:**
  - Wrap the model in a Flask or FastAPI service and deploy to Vercel or Hugging Face Spaces.

---

## 8. Dependencies

- transformers
- datasets
- sentencepiece
- scikit-learn
- nltk
- torch
- bert\_score

---

## 9. Potential Improvements

- Incorporate domain-specific data for fine-tuning
- Use larger context-aware models (e.g., mBART or T5 multilingual)
- Introduce post-processing grammar correction or reranking
- Utilize semantic-aware metrics (e.g., BERTScore) alongside BLEU
- 

## 10.RUN

**Translation English**

Enter English sentence to translate to Arabic

input\_text

I have to go home.

◆ Predict Translation

يجب أن أذهب إلى البيت

Clear

Submit

Flag

**Translation English**

Enter English sentence to translate to Arabic

input\_text

My tie is orange.

◆ Predict Translation

ربطة عنقي برتقالية

Clear

Submit

Flag

**Translation English**

Enter English sentence to translate to Arabic

input\_text

Please don't cry.

◆ Predict Translation

من فضلك لا تبكي

Clear

Submit

Flag

## 11. Conclusion

This project successfully demonstrates the application of a pre-trained MarianMT model for Arabic-English translation. Through fine-tuning and bertscore evaluation, the model achieved good performance on a general dataset, with opportunities for further enhancement in specialized applications.

-